

Load Balancing for Traffic in Networks

(with Dóra Erdős)


Feiyu Shi
4/29/2014

Recap

- + Problem:
- + Given a set of k nodes, partition it into **two** groups such that the difference of **total number of shortest paths covered by each group** is minimal.
- + Steps: node centrality \rightarrow group centrality \rightarrow load balancing
- + ~~Original setting: in Undirected Graph~~
- + Current setting: in Directed Acyclic Graph (DAG)

Example Graph

- + Given a group of nodes: A, B, C, D, E, F, divide this group into two groups, such that the difference of the group centralities of these two groups are minimized.



Node Centrality

- + Prefix: $PREFIX(v) = \#PATHS(S, v) = \sum_{s \in S} \#PATHS(s, v)$
- + Suffix: $SUFFIX(v) = \#PATHS(v, T) = \sum_{t \in T} \#PATHS(v, t)$
- + Impact: $I(v) = PREFIX(v) \times SUFFIX(v)$
- + From general paths to shortest paths: check if parents of v are on a shortest path. If $d(s, v) - d(s, x) = 1$, then x is in $Pi'(v)$. Only includes parents from $Pi'(v)$ in the computation of prefix and plist (suffix).

Compute Prefix

- + Observation: every path from source s to a node v has to go through one of v 's parents.
- + Recurrence relation:
- + Base: $PREFIX(s) = 1$
- + Induction: $PREFIX(v) = \sum_{x \in Pi(v)} PREFIX(x)$
- + Order: topological order determines ancestors of v .
- + Thus, $PREFIX(v) = \# PATHS(S, v) = \sum_{x \in Pi(v)} \# PATHS(S, x)$

Compute Suffix

- + Plist_v : contains for every ancestor y of v the number of paths that go from y to v .
- + Another recurrence relation:
- + Base: $PLIST_y[y] = 1$
- + Induction: $PLIST_v[y] = \sum_{x \in \text{Pi}(v)} PLIST_x[y]$
- + Finally: $SUFFIX(v) = \sum_{t \in T} PLIST_t(v)$

Group Centrality

- + Suppose current set A 's group centrality is $C(A)$. After we add another node v to set A , we get the new group centrality $C'(A)$.
- + Define the conditional impact: $I_A(v) = C'(A) - C(A) \leq I(v)$.
- + If $A = \text{empty set}$, $I_{\emptyset}(v) = I(v)$
- + It's the number of shortest paths that were not covered by A .
- + Set $\text{Prefix}(a)$ to zero;
- + Set $\text{Plist}_v(a)$ to zero.

Load Balancing: Problem

- + Let $G(V, E)$ be a DAG, where $|V| = n$. Given a group of nodes, $V_0 \subseteq V$ where $|V_0| = k$. Find **an assignment** of two groups V_1 and V_2 , where $V_1 \cap V_2 = \emptyset$ and $V_1 \cup V_2 = V_0$, such that the difference of group centralities of V_1 and V_2 is minimized.
- + Recall: brute-force algorithm and greedy algorithm

Baseline: Brute-force Algorithm

- + Enumerate all possible assignments: 2^k combinations.
- + The evaluate the group centralities of each possibility.
- + Find the solution having the minimal group centrality difference.
- + Slowest, but can find **all the best** assignments.

Implementation of Brute-force Alg.

- + First, find the power set of the set of k nodes.
- + For example for set of $\{A, B, C\}$, its power set is $\{\emptyset, \{A\}, \{B\}, \{C\}, \{A, B\}, \{A, C\}, \{B, C\}, \{A, B, C\}\}$.
- + Second, compute the group centrality of each set in the power set.
- + Third, compute the difference of one set and its complement set. Find the minimum of the differences.
- + Finally, find all assignment of sets that equals the min diff.

Greedy Algorithm

- + Inspired by the concept of conditional impact.
- + If we move node v from G_1 to G_2 , then the new centralities of two groups $C_1' = C_1 - I_{G_1}(v)$, $C_2' = C_2 + I_{G_2}(v)$.
- + Then new difference: $C_1' - C_2' = (C_1 - C_2) - (I_{G_1}(v) + I_{G_2}(v))$.
- + Heuristic: choosing the node whose new difference is 'absolutely' closest to the previous difference.

Implementation of Greedy Algorithm

- + Input: $G(V, E)$, V_0
- + Output: V_1 or V_2
- + 1. $V_1 = V_0$, $V_2 = \emptyset$
- + 2. Compute $\text{diff} = C_{sp}(V_1) - C_{sp}(V_2)$.
- + 3. Compute $\text{diff}' = C'_{sp}(V_1) - C'_{sp}(V_2)$. Find v_i whose $|\text{diff}'|$ is minimal (there may not be only one v_i , but randomly choose only one).
- + 4. Remove the node from V_1 and add it to V_2 .
- + if $|\text{diff}'| \leq |\text{diff}|$ and V_1 is not empty then
 - + 5. recursively solve the problem with new V_1 , V_2 .
- + else
 - + 6. return old V_1 or V_2 .
- + end if.

Improvement: Greedy Search Alg.

- + Based on greedy algorithm.
- + Can find more than one solution if there exists many solutions. But it's not guaranteed that it can find all the solutions. (Remember brute-force alg. can find all the solutions.) The order to pick nodes matters.
- + New: Backtracking search. Try all possibilities. Then choose the best combination.
- + New: Pruning. Maintaining a list 'visited'.

Implementation of Greedy Search Alg.

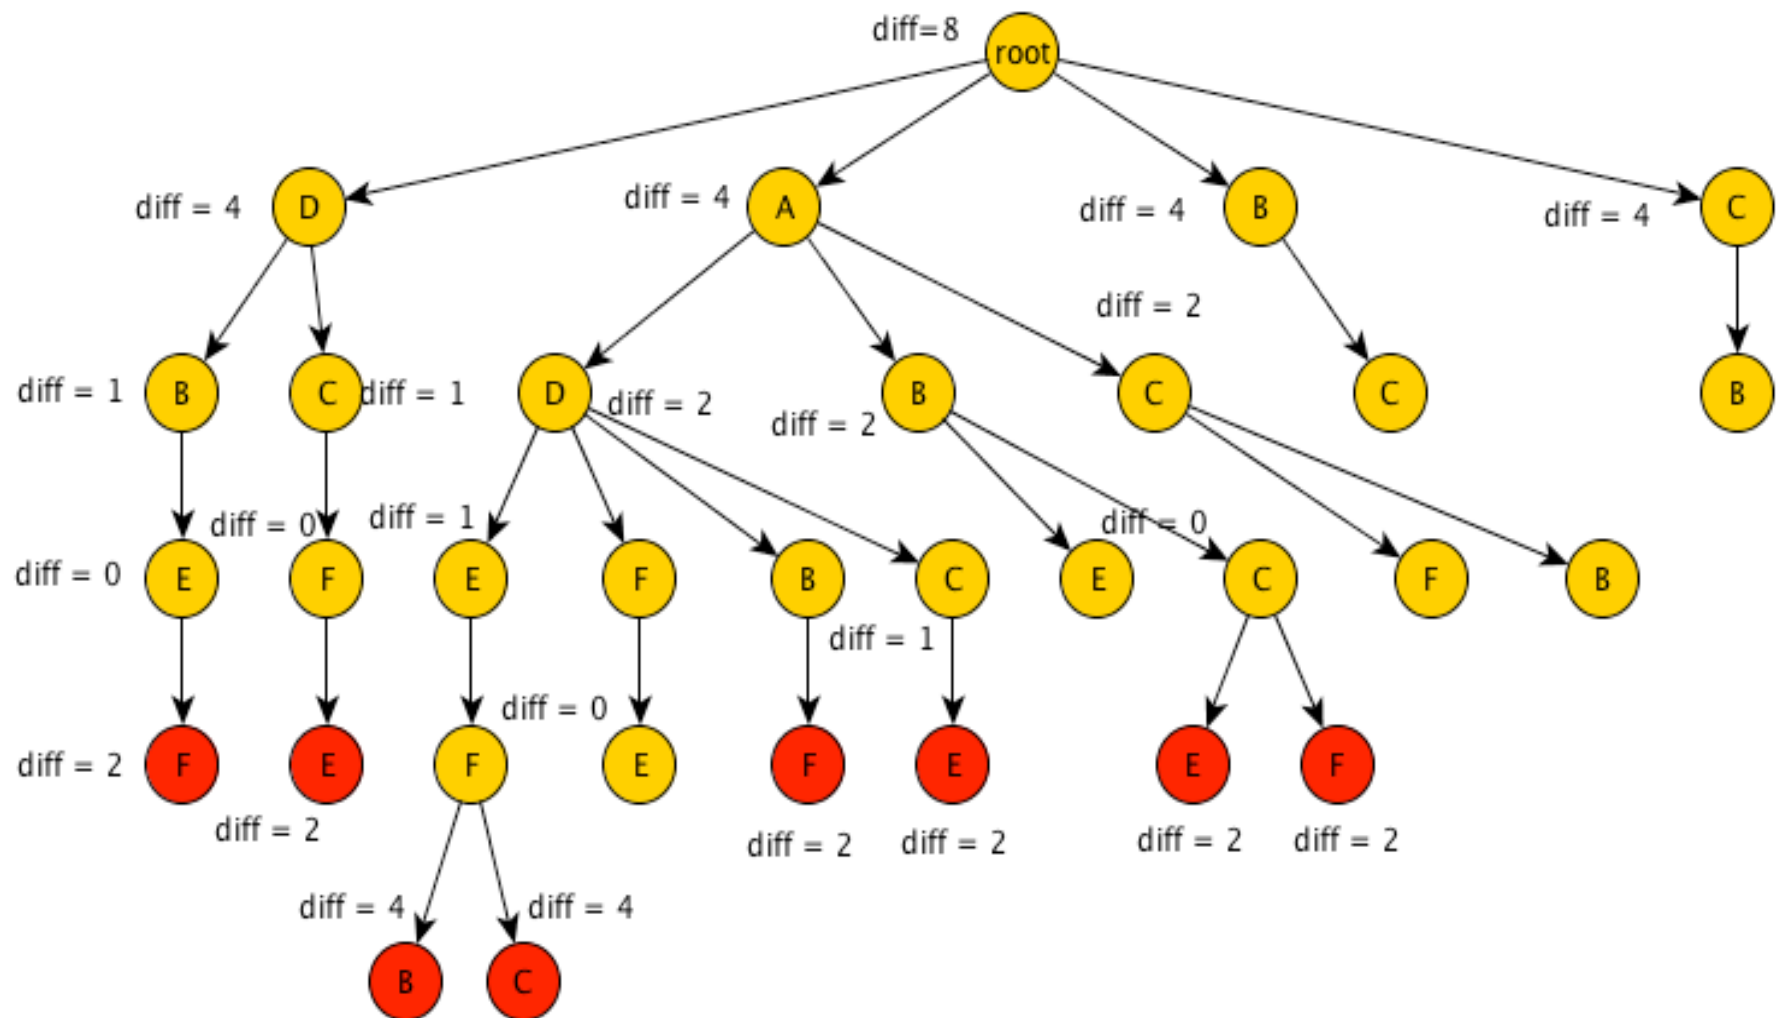
- + Input: $G(V,E)$, V_0
- + Output: A list of V_1 or V_2
- + 1. $V_1 = V_0$, $V_2 = \emptyset$
- + 2. Compute $\text{diff} = C_{sp}(V_1) - C_{sp}(V_2)$.
- + 3. Compute $\text{diff}' = C'_{sp}(V_1) - C'_{sp}(V_2)$. Find **all** v_i whose $|\text{diff}'|$ is minimal

- + Return a list of assignments by choosing the best ones with $\min |\text{diff}|$.

Greedy Search Alg. (Cont.)

- + for each v_i do
 - + 4. Remove v_i from V_1 and add it to V_2 .
 - + If neither new V_1 nor new V_2 is in the list 'visited' then
 - + 5. Compute the new diff'.
 - + If $|\text{diff}'| < |\text{diff}|$ and V_1 is not empty then
 - + If $\text{diff}' = 0$ then add new V_1, V_2 to list 'assign' and 'visited'.
end if
 - + 6. recursively solve the problem with new V_1, V_2 .
 - + else
 - + 7. add old V_1 or V_2 to list 'assign' and 'visited'.
 - + end if
 - + end if
 - + 8. Return a list of assignments from 'assign' with min difference
- + end for

Search Graph of Greedy Search Alg.



Full Search Algorithm

- + Similar to greedy search algorithm, but need to compute group centralities for each node and relax the constraint:
 - + Instead of using greedy heuristic, we loop over **all possible nodes**.
- + Give all possible solutions, the same as brute-force method.
- + Spends more time (even more than brute-force).

Time Complexity Analysis

- + Recall: the time complexity for computing group centrality is $O(k |E| \Delta)$. Node centrality is $O(|E| \Delta)$.
- + Brute-force: $O(2^k k |E| \Delta)$.
- + Greedy: k level. 1 node / level. So $O(k^2 |E| \Delta)$.
- + Full search: worst case $k!(1 + 1/2! + 1/3! + \dots + 1/(k-1)!) \approx k! e$. So $O(k! k |E| \Delta)$.
- + Greedy search: somewhere between $O(k^2 |E| \Delta)$ and $O(k! |E| \Delta)$, and usually $< O(2^k k |E| \Delta)$ (from test result).



Experiments

Dataset Preparation, Validation, Time complexity

Dataset (Thanks to Dóra)

- + Synthetic DAG dataset:
- + Three parameters: n , d , l . n is the number of nodes in the graph, d is a density parameter and l is a limit.
- + Node are labeled as integers. For every node i to node j , we try to generate an edge with probability d if $i < j \leq i+l$. $d \cdot l$ = expected degree of every node. $|E| \approx d \cdot l \cdot n$.
- + We generated
 - + Dataset1: 10 graphs with varying d from 0.1 to 1.0 ($n=200$, $l=10$).
 - + Dataset2: 10 graphs with varying n from 100 to 1000 (not used).
 - + Data: 1 DAG, $n = 50$, $d = 0.4$, $l = 5$.
- + We compare performance of three algorithms on these datasets.

Dataset (cont.)

- + Dataset1: Choose first 10 nodes as sources, last 10 nodes as destinations.
- + DAG 50_0.4_5: first 5 nodes as sources, last 5 nodes as destinations.
- + Choice of group:
 - + Dataset1: 5 nodes. For graph with n nodes, choose $0.4n$, $0.45n$, $0.5n$, $0.55n$, $0.6n$.
 - + DAG 50_0.4_5: 10 groups with number of nodes from 4 to 13. Smaller group \subseteq larger group.

Tests

- + Two aspects: validation and time complexity.
 - + Validation: test if the algorithm returns (all) the correct answer.
 - + Time complexity: measure the running time for each alg. in ms.
- + Two set of experiments:
 - + On dataset1: focus on validation
 - + On DAG 50_0.4_5: focus on time complexity (k).

Dataset1: Validation


Test	1	2	3	4	5	6	7	8	9	10
num.	16	16	16	8	4	2	1	1	1	1
G	Y	Y	Y	Y	Y	Y	Y	Y	Y	Y
GS #	16	16	16	7	3	1	1	1	1	1
GS	Y	Y	Y	Y	Y	Y	Y	Y	Y	Y
R %	100	100	100	87.5	75	50	100	100	100	100
FS #	16	16	16	8	4	2	1	1	1	1
FS	Y	Y	Y	Y	Y	Y	Y	Y	Y	Y
R %	100	100	100	100	100	100	100	100	100	100

1. num.: total number of solutions (from baseline method).
2. G: greedy validation.
3. GS #: number of solutions from greedy search method.
4. GS: greedy search validation.
5. R %: retrieval rate of greedy search method; $GS \# / num.$

Dataset1: Validation (cont.)

- + Special cases:
 - + 1. 0 vs. 0 (test 1)
 - + 2. largest vs. 0 (test 2, 3)
- + Sample test results:
 - + 227568 vs. 19836
 - + 32148083 vs. 43532838

Dataset1: Time Complexity




DAG 50_0.4_5: Validation

k	4	5	6	7	8	9	10	11	12	13
num.	1	1	2	4	4	8	8	8	8	8
G	Y	Y	Y	Y	Y	Y	N	Y	Y	N
GS #	1	1	1	3	3	7	0	7	7	0
GS	Y	Y	Y	Y	Y	Y	N	Y	Y	N
R %	100	100	50	75	75	87.5	0	87.5	87.5	0
FS #	1	1	2	4	4	8	8	8	8	8
FS	Y	Y	Y	Y	Y	Y	Y	Y	Y	Y
R %	100	100	100	100	100	100	100	100	100	100


Note:

Impacts from greedy-like alg (k = 10): 751 vs. 733. Baseline: 711 vs. 700.

DAG 50_0.4_5: Time Complexity



DAG 50_0.4_5: Time Complexity



Conclusions

- + Greedy algorithm is good enough (90% chance) **sometimes**.
- + Greedy search algorithm works **as good as** greedy algorithm. But it might be suitable when there is a need for more than one possible solutions.
- + Brute-force and full search algorithms are able to get **all** possible solutions.
- + Future study may be applying these algorithms to undirected graphs.

Questions?